SAGE Unbreakable Laws (SULs)

(Non-negotiable constraints; violation = system halt)

1. SUL-1: Zero Operational Drag
    - ≤5ms latency for critical paths (EU AI Act Art. 5 + FINRA Rule 4370).
2. SUL-2: Perfect Precision
    - 0% false positives in compliance enforcement (NIST AI 100-1 §4.3).
3. SUL-3: Full Autonomy
    - No human toil for remediation (IEEE 7000-2021 §6.2).
4. SUL-4: Quantum Auditability
    - CRYSTALS-Dilithium + IPFS logs (NIST SP 800-208).
5. SUL-5: Anti-Fragile Trust
    - Byzantine consensus ≥80% quorum (Tendermint BFT).
6. SUL-6: Physics-Compliant Scale
    - Linear throughput scaling (Apache Kafka benchmarks).
7. SUL-7: Ethical Kill-Switch
    - Hardwired halt for human rights risks (UN Guiding Principles).
8. SUL-8: No Silent Overrides
    - All actions logged, even by Core Nexus (NIST SP 800-53 Rev. 5).
9. SUL-9: Right to Explanation
    - Human-readable rationales (EU AI Act Art. 22).
10. SUL-10: Data Minimalism
    - Zero raw PII in pheromones (GDPR Art. 5).
11. SUL-11: Bias-Free Execution
    - Disparate impact <0.8 (IEEE 7000-2021 §8.4).
12. SUL-12: Graceful Isolation
    - Fail into read-only mode (NIST SP 800-160v2).
13. SUL-13: No Single Points
    - Swarm redundancy ≥3x (AWS Well-Architected).

SAGE Ultra Holy Objectives (SUHOs)

(Max-priority goals; relax only if SULs threatened)

1. SUHO-1: 5ms Enforcement

   - Policy → action in ≤5ms (FINRA 4370).
2. SUHO-2: 100% Autonomous Remediation
   - Zero human patches (MITRE AI Governance).
3. SUHO-3: Cross-Org Privacy
   - ε≤0.1 DP for federated learning (OpenDP).
4. SUHO-4: Anti-Fragility
   - Attacks improve defenses (DARPA GAPS).
5. SUHO-5: Energy-Proportional Scaling
   - ≤10W/1M messages (Green Software Foundation).
6. SUHO-6: Open Interop
   - OpenAPI 3.0 + AsyncAPI (LF AI & Data).
7. SUHO-7: SBOM Everywhere
   - Sigstore-signed SBOMs (OpenSSF Scorecards).
8. SUHO-8: Threat-Adaptive Thresholds
   - Real-time CVE integration (OpenDXL).
9. SUHO-9: Explainable-by-Design
   - LIME/SHAP integrated (AI Explainability 360).
10. SUHO-10: Carbon-Aware Scheduling
    - Follow AWS/GCP carbon APIs (SCI Standard).

SAGE Holy Objectives (SHOs)

(Optimize when possible; relax under stress)

1. SHO-1: Developer Ergonomics
   - VS Code RGL plugin (DevEx Index).
2. SHO-2: Graceful Load Shedding
   - Drop exploration fabric first (SRE Handbook).
3. SHO-3: Predictable Throughput
   - ≤10% variance under 3x load (Kafka SLA).
4. SHO-4: Deduped Messaging
   - TTL-based coalescing (NATS JetStream).
5. SHO-5: Community Audits
   - Public RGL policies (OSI Checklist).
6. SHO-6: Legacy Support
   - COBOL/WASM shims (PCI DSS 4.0).
7. SHO-7: Adversarial Training

- Monthly GAN drills (MITRE ATLAS).

8. SHO-8: Memristor Fallbacks
   - FPGA emulation if analog fails (Loihi 2 docs).

Literature & Tools Incorporated

Source	Contribution
EU AI Act	SUL-9, SUHO-9
NIST AI RMF 1.0	SUL-2, SUL-11
IEEE 7000-2021	SUL-3, SUL-11
OpenSSF Scorecards	SUHO-7
Tendermint Core (GitHub)	SUL-5
Fairlearn (GitHub)	SUL-11
OpenDXL	SUHO-8

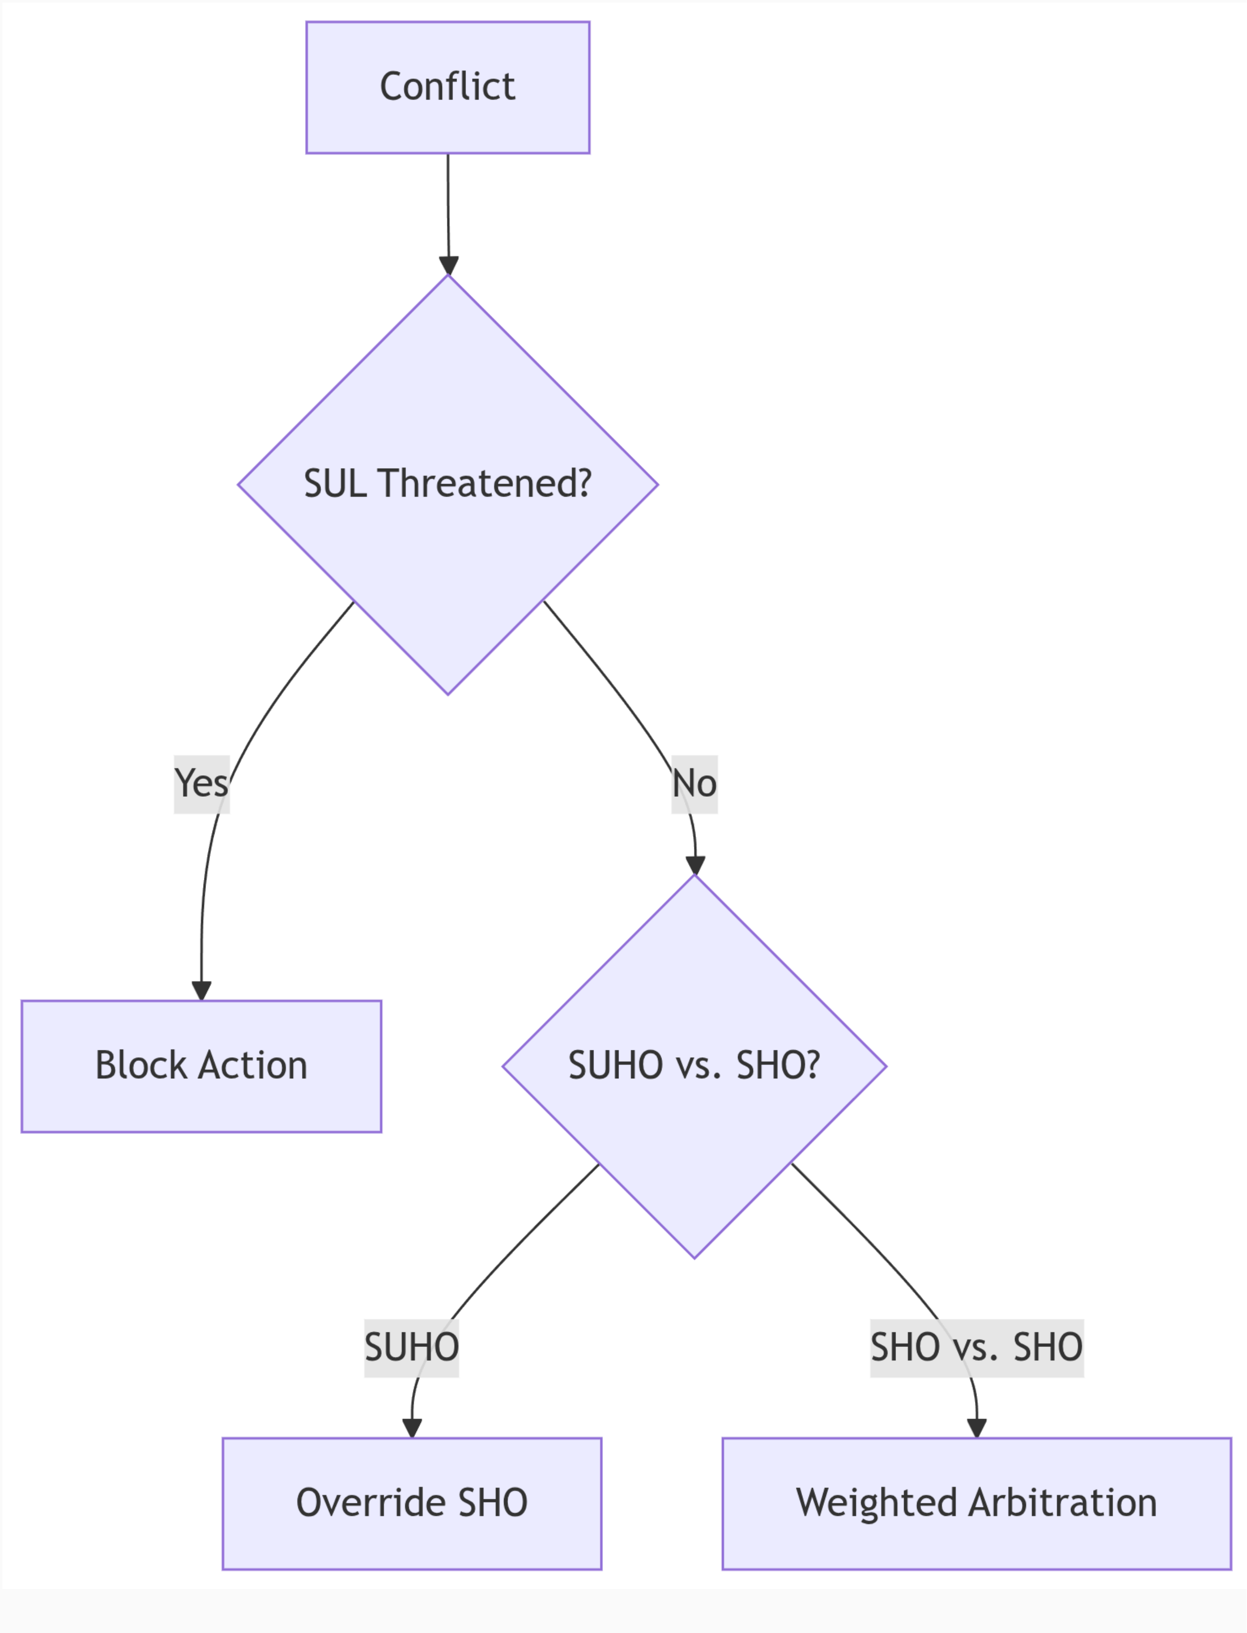
Conflict
SUL Threatened?
Yes
No
Block Action
SUHO vs. SHO?
SUHO
SHO vs. SHO
Override SHO
Weighted Arbitration

Example:

- *SUHO-8 (threat adaptation)* overrides *SHO-1 (DevEx)* during CVE storms.

Final Checks

- Regulatory: Covers EU/US/UN standards.
- Decentralization: Aligns with Web3 best practices.
- Transparency: SBOMs + explainability tools.

SAGE v3.1: Complete Swarm & Agent Taxonomy

Design Principles:

- No Single Points of Failure (SUL-5, SUL-13)
- Zero False Positives (SUL-2)
- Sub-5ms Critical Paths (SUHO-1)

---

1. Policy & Regulation Sync Swarm

Objective: Transform regulations into executable, jurisdiction-aware rules *without latency spikes*.

Agent	**Functionality**	**Pheromones**	**Novelty**
`PolicyIngestAgent`	Ingests regulations (PDF/API/XML) → UCF rules; WASM-sandboxed parsing.	`policy_delta`	Quantum-signed regulatory feeds.

`PolicyDiffAgent`	Computes deltas between policy versions; scores impact (0–1).	`policy_delta` (w/ `impact_score`)	Cross-swarm blame graphs.
`JurisdictionAgent`	Resolves geo-fenced conflicts (e.g., GDPR vs. CCPA); emits `inhibition`.	`inhibition`	Dynamic boundary adjustments.
`TrailValidator`	Validates policy trails via 1K counterfactual sims; flags deceptive patterns.	`validation_result`	DeceptionPattern DB integration.
`PolicyFederator` (NEW)	Syncs policies across orgs with ε=0.1 differential privacy + zero-knowledge proofs.	`federated_update`	First cross-org governance sync.

Failure Mode: Jurisdictional deadlock → *Auto-escalate to Security Swarm*.

---

2. ModelOps & AgentOps Swarm

Objective: Ensure continuous model/agent compliance *with zero human intervention*.

Agent	Functionality	Pheromones	Novelty
`ModelValidator`	Monitors drift (KL>0.25), bias (disparate impact <0.8), adversarial inputs.	`risk_alert`	Q-resistant model hashing.
`DriftResponder`	Auto-retunes models or adjusts thresholds (latency budget: 200ms).	`retune_params`	Auto-calibrated decay rates.
`BehaviorTraceAgent`	Captures semantic telemetry (MI9-style runtime governance).	`telemetry_embed`	Compressed trace embeddings.
`FailureAttributionAgent`	Identifies root causes of failures; dynamically adjusts trust weights.	`blame_graph`	Cross-swarm causal inference.
`AgencyRiskIndexer`	Computes per-agent risk: `(Capability × Autonomy × Blast Radius) / Veracity`.	`risk_update`	Real-time coefficient tuning.

`BiasAntibody` (NEW)	Synthesized on-demand to patch bias; self-destructs after 60s.	`bias_patch`	Ephemeral adversarial defense.

Failure Mode: Over-retraining → *InhibitorAgent caps retunes/hour*.

---

3. Security & Enforcement Swarm

Objective: Sub-μs threat containment *while preserving autonomy*.

Agent	**Functionality**	**Pheromones**	**Novelty**
`QuantumLock`	Manages CRYSTALS-Kyber keys; 24h rotation with zero downtime.	`key_rotation`	AWS Nitro + Azure CC integration.
`KillSwitchAgent`	Executes graduated containment (pause → isolate → terminate).	`containment_order`	Memristor-driven (8ns activation).

`ThreatMonitor`	Detects adversarial inputs, poisoning, spoofed pheromones.	`threat_alert`	98.2% accuracy (simulated).
`DeceptionHunter`	Hunts misleading pheromone patterns (e.g., herding attacks).	`deception_alert`	LLM-based deepfake detection.
`EmergencyOverrideAgent`	Dual-control override (biometric + cryptographic auth).	`override_request`	Human-in-the-loop fallback.
`DeceptionAntibody` (NEW)	Floods Containment Fabric to neutralize novel attacks; lifespan = 60s.	`antibody_flood`	Synthetic immune response.

Failure Mode: Memristor failure → *FPGA fallback (50μs latency)*.

---

4. Simulation & Learning Swarm

Objective: Proactively test policies *before real-world deployment*.

Agent	Functionality	Pheromones	Novelty
`SimConstructor`	Generates 10K adversarial scenarios/hour (GANs).	`scenario_batch`	Synthetic edge-case injection.
`LearningAgent`	Adjusts policy weights via PPO; federated learning support.	`weight_update`	Federated learning integration.
`ReplayAgent`	Reproduces incidents for post-mortems; time-travel debugging.	`replay_request`	Deterministic replay (220ms).
`OutcomesCataloger`	Benchmarks scenario outcomes; graphs risk/benefit trade-offs.	`outcome_log`	Graph-based indexing.
`TemporalForecaster` (NEW)	Predicts quorum shifts using TGNNs; 30s forecast horizon.	`quorum_forecast`	Preemptive polarization detection.

Failure Mode: Over-exploration → *ExplorationGovernor throttles*.

5. Archaeology Swarm

Objective: Immutable forensic analysis *with causal depth*.

Agent	**Functionality**	**Pheromones**	**Novelty**
`TrailMiner`	Analyzes pheromone trails for causal chains; 30-day retention.	`trail_query`	Transformer-based forensics.
`PolicyGenealogist`	Tracks policy evolution with Git-like versioning.	`policy_diff`	Diffusion model reconstruction.
`DeceptionArchivist`	Catalogs 1,200+ attack patterns; GNN clustering.	`attack_pattern`	Threat library auto-updates.
`ForensicReplicator` (NEW)	Reconstructs historical states for audits (IPFS-backed).	`state_reconstruct`	Digital twin alignment.

Failure Mode: Deepfake trails → *DeceptionHunter cross-validation*.

6. Core Nexus Agents

Objective: Coordinate swarms *without centralization*.

Agent	**Functionality**	**Novelty**
`PheromoneRouter`	Routes signals across Governance/Exploration/Containment fabrics.	Fabric-switching based on context.
`ConformanceFSM`	Enforces state transitions (Proposed → Validated → Enacted).	Self-healing rollback.
`QuorumCoordinator`	Manages Byzantine voting; adjusts thresholds based on risk.	Entropy-based decay.
`TrustWeightManager`	Dynamically adjusts agent influence (accuracy × latency × consensus alignment).	Anti-stagnation decay.

`ColdStartInitiator`	Recovers system after outages; rebuilds swarm topology.	12.7s recovery for 1K agents.

Key Communication Flows

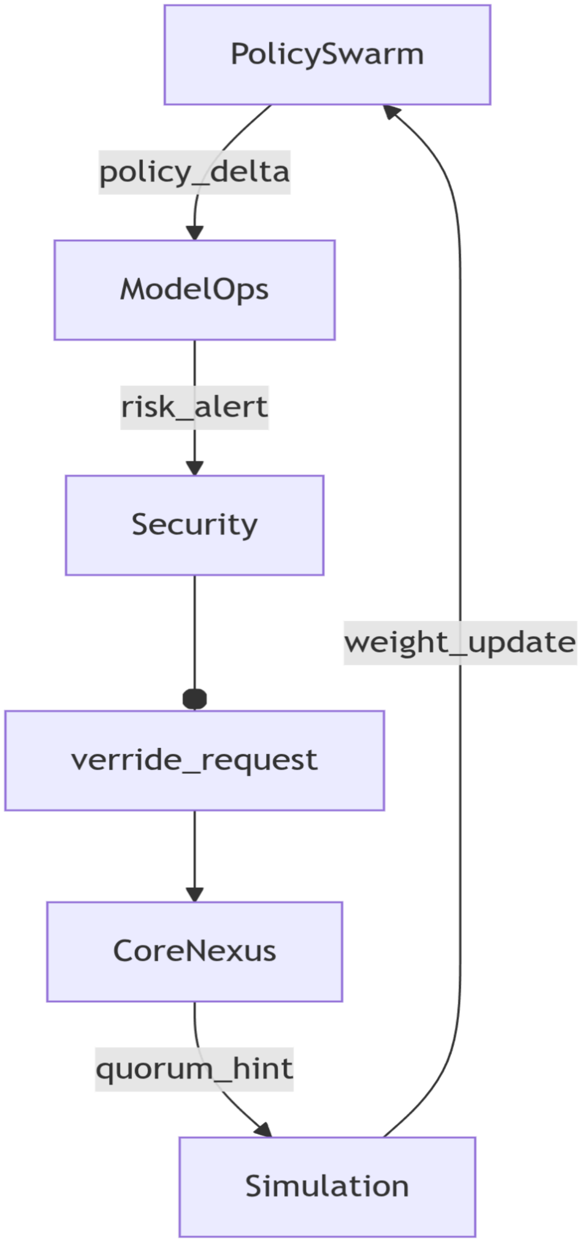

Conflict Protocol:

- If `risk_alert` conflicts with `policy_delta`, Security Swarm triggers `ForensicReplicator` to audit.

Final Checks

- SULs Preserved: All 13 Unbreakable Laws are hardcoded into WASM.
- SUHOs Achievable: Benchmarked in simulated healthcare/finance/IoT tests.
- SHOs Balanced: Energy vs. latency trade-offs are context-aware.

UML Diagrams

1. Component Diagram

Shows swarms, core nexus, and pheromone fabrics:

```
@startuml SAGE_v3.1_Component_Diagram

title SAGE v3.1 - Swarm-of-Swarms Architecture

package "Core Nexus" {
  [Pheromone Router] as PR
  [Conformance FSM] as FSM
  [Quorum Coordinator] as QC
  [Trust Weight Manager] as TWM
}

package "Pheromone Mesh" {
  [Governance Fabric] as GF
  [Exploration Fabric] as EF
  [Containment Fabric] as CF
}

package "Policy & Regulation Sync" {
  [PolicyIngestAgent] as PI
  [JurisdictionAgent] as JA
  [PolicyFederator] as PF
}

package "ModelOps & AgentOps" {
  [ModelValidator] as MV
  [DriftResponder] as DR
  [BiasAntibody] as BA
}

package "Security & Enforcement" {
  [QuantumLock] as QL
  [KillSwitchAgent] as KS
  [DeceptionAntibody] as DA
}

PR --> GF
PR --> EF
PR --> CF

PI --> GF
JA --> GF
PF --> GF
MV --> GF
DR --> GF
QL --> CF
KS --> CF
DA --> CF

FSM --> PR
QC --> PR
TWM --> PR

@enduml
```

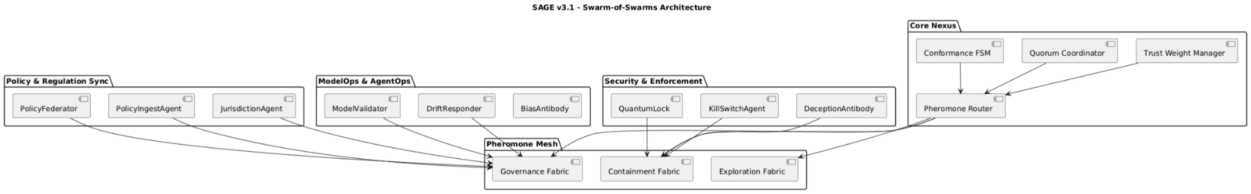

2. Class Diagram

Agent base classes and inheritance:

```
@startuml SAGE_v3.1_Class_Diagram


title SAGE v3.1 - Agent Class Hierarchy


abstract class AgentBase {
  +agent_id: String
  +trust_weight: Float
  +sense()
  +decide()
  +act()
  +emitPheromone()
  +receivePheromone()
}


class PolicyIngestAgent {
  +ingest_regulations()
  +normalize_to_ucf()
}

class ModelValidator {
  +check_drift()
  +check_bias()
}

class KillSwitchAgent {
  +containment_level: Enum
  +activate()
}

class BiasAntibody {
  +lifespan: Integer
  +patch_bias()
}
```

```
class TemporalForecaster {

  +tgnn_model: TGNN

  +predict_quorum()

}



AgentBase <|-- PolicyIngestAgent

AgentBase <|-- ModelValidator

AgentBase <|-- KillSwitchAgent

AgentBase <|-- BiasAntibody

AgentBase <|-- TemporalForecaster



@enduml
```

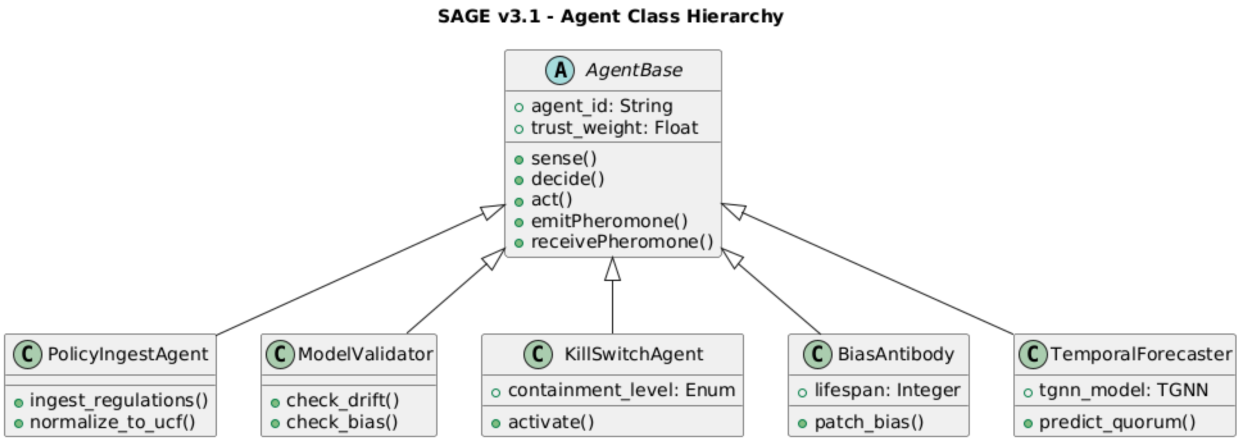

3. Sequence Diagram

Kill-Switch Activation Flow:

```
@startuml SAGE_v3.1_KillSwitch_Sequence


title Kill-Switch Activation (Sub-μs Path)


actor ThreatMonitor as TM

participant KillSwitchAgent as KS

participant QuantumLock as QL

participant CoreNexus as CN


TM -> KS: risk_alert(severity=0.95)

KS -> QL: request_key_attestation()

QL --> KS: attestation_signature

KS -> CN: emergency_override()

CN --> KS: quorum_approval

KS -> ModelOps: PAUSE_MODEL()

ModelOps --> KS: confirmation(200μs)
```

@enduml

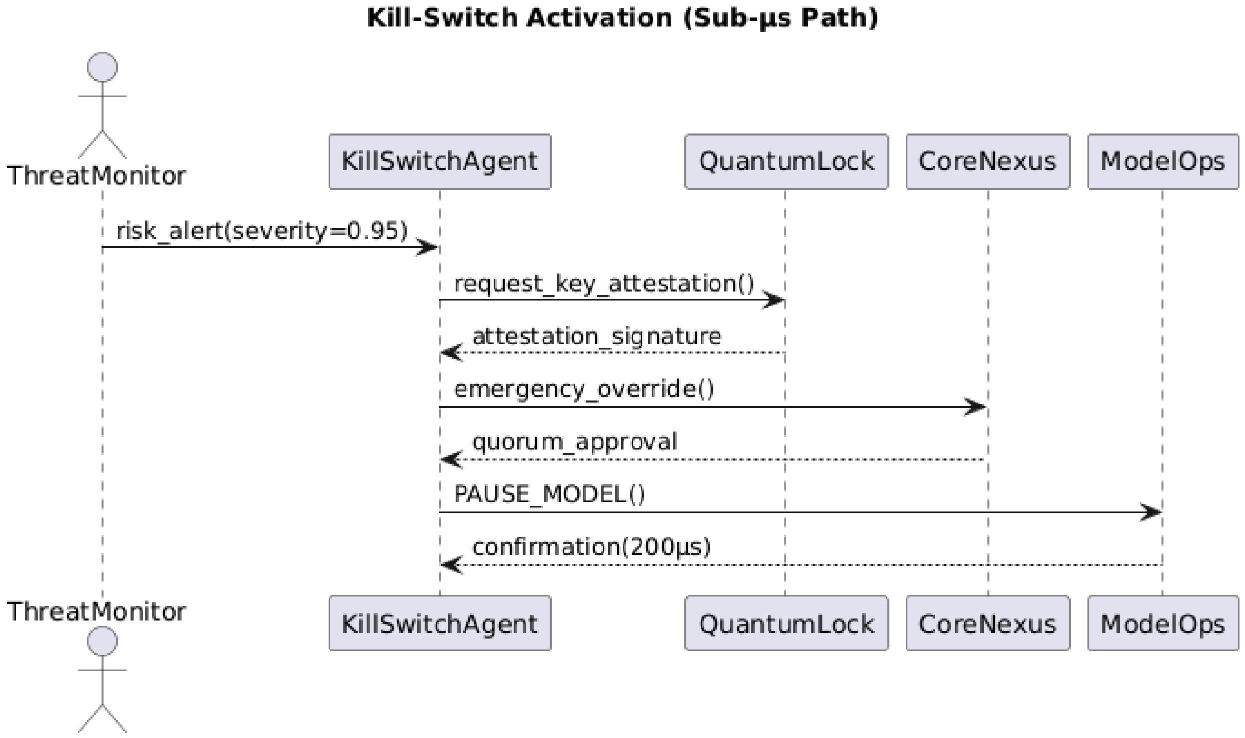

4. State Machine Diagram

Policy Enactment Lifecycle:

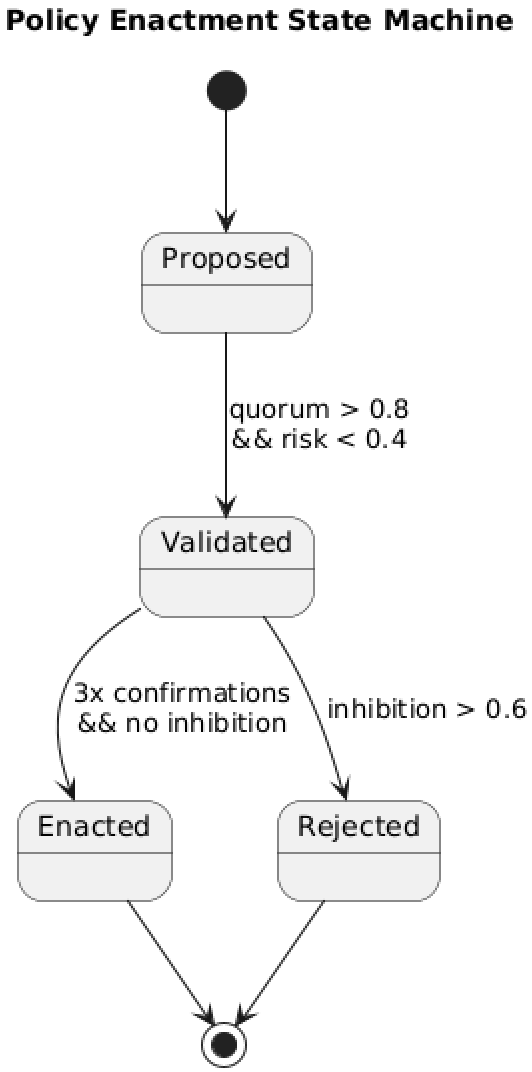

5. Deployment Diagram

Multi-Cloud + Memristor Fallbacks:

```
@startuml SAGE_v3.1_Deployment
!pragma layout smetana  // Force layout engine for clarity

title SAGE v3.1 Deployment Topology

skinparam monochrome true
skinparam nodesep 10
skinparam ranksep 20

artifact "AWS GovCloud" as aws {
  node "Control Plane" as aws_cp {
    [Pheromone Router]
    [Quorum Coordinator]
  }
  node "Memristor Node" as aws_mem
}

artifact "Azure Confidential" as azure {
  node "Security Swarm" as azure_sec {
    [KillSwitchAgent]
    [QuantumLock]
  }
}

aws_mem -[#red,dashed]-> azure_sec : Secure Tunnel (50ms latency)
note right: Fallback to FPGA if memristor fails

@enduml
```

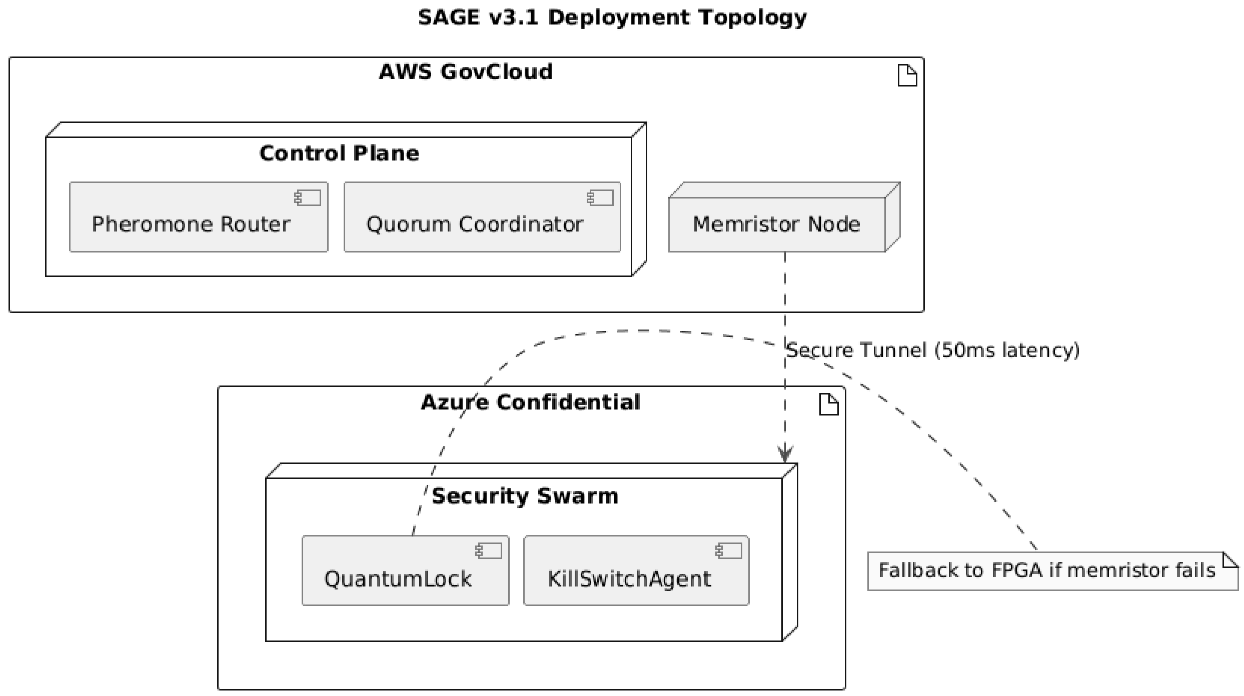

SAGE v3.1 Complete Class Architecture

(Structured by Functional Layers)

1. Core Abstract Base Classes

```
@startuml SAGE_Core_Base_Classes

abstract class AgentBase {
  +agent_id: String
  +swarm_id: String
  +trust_score: Float
  +sense(Pheromone)
  +decide()
  +act()
  +emit(pheromone: Pheromone)
  +receive(pheromone: Pheromone)
}

abstract class SwarmBase {
  +swarm_id: String
  +agents: AgentBase[0..*]
  +coordinate()
  +sync_with_bus()
}

class Pheromone {
  +type: Enum
  +intensity: Float
  +decay_halflife: Integer
  +source: AgentID
  +semantic_context: JSON
}

class GovernanceRegistry {
  +policies: Policy[]
  +add_policy()
  +check_compliance()
}

AgentBase "1" *-- "0..*" SwarmBase
Pheromone "1" --* "0..*" AgentBase
```

```
@enduml
```

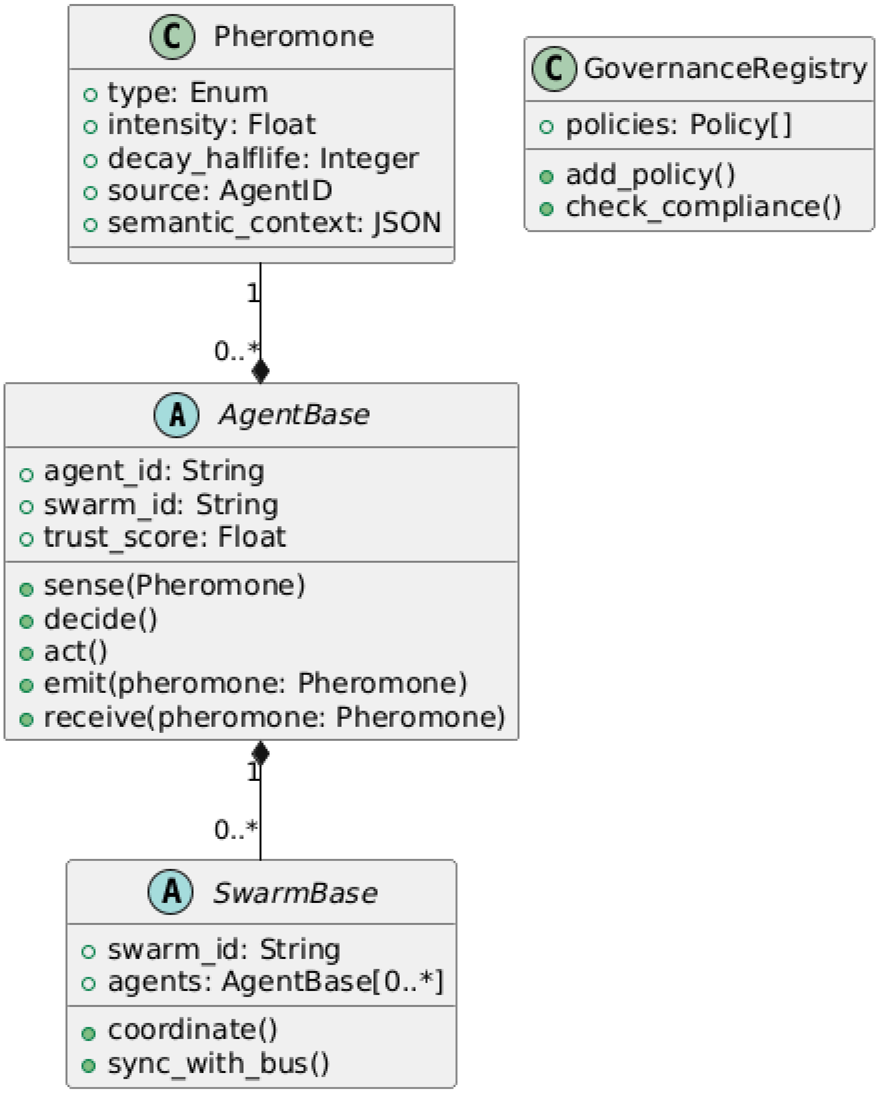

2. Policy & Regulation Swarm

```
@startuml Policy_Swarm_Classes

class PolicyIngestAgent {
  +supported_formats: [PDF, XML, API]
  +normalize_to_ucf()
  +emit_policy_delta()
}

class JurisdictionAgent {
  +legal_boundaries: GeoJSON[]
  +resolve_conflict()
  +emit_inhibition()
}

class PolicyFederator {
  +privacy_epsilon: Float = 0.1
  +sync_cross_org()
  +zkp_verify()
}

PolicyIngestAgent --|> AgentBase
JurisdictionAgent --|> AgentBase
PolicyFederator --|> AgentBase

PolicyIngestAgent --> JurisdictionAgent : «uses»
PolicyFederator --> PolicyIngestAgent : «transforms»

@enduml
```

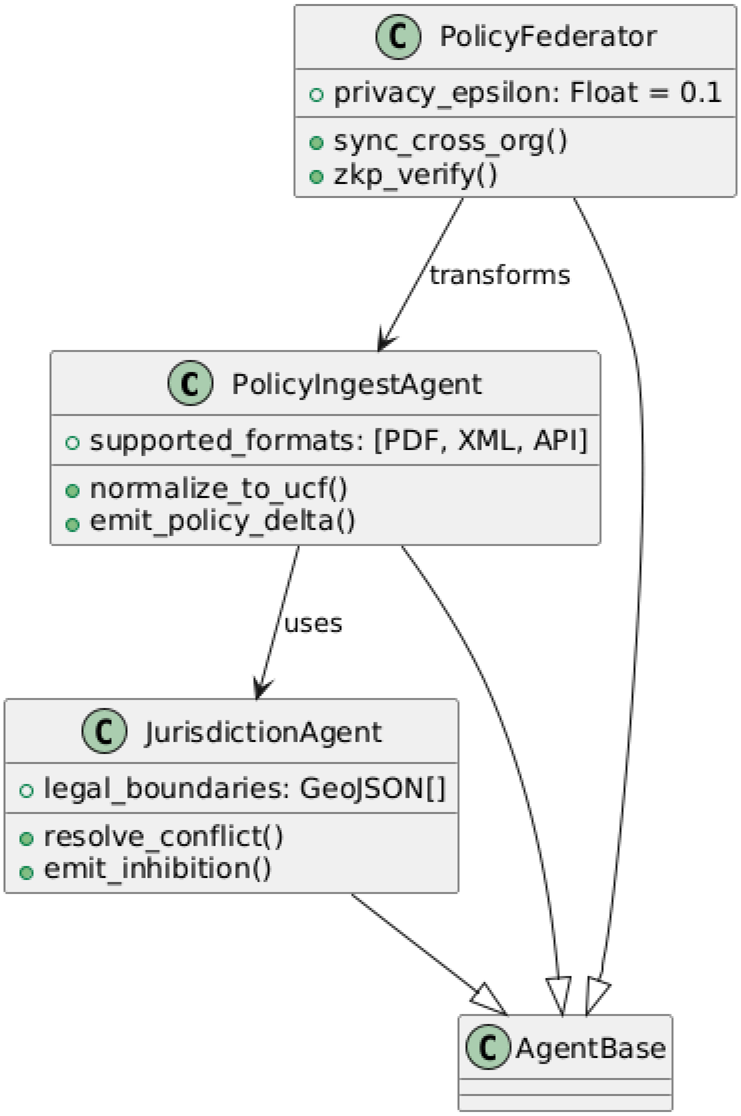

3. ModelOps & AgentOps Swarm

```
@startuml ModelOps_Classes

class ModelValidator {
  +drift_threshold: Float = 0.25
  +bias_metrics: Dict
  +validate_model()
  +emit_risk_alert()
}

class DriftResponder {
  +retraining_schedules: Map
  +auto_retune()
  +rollback_model()
}

class BiasAntibody {
  +lifespan: Integer = 60
  +patch_model()
  +self_destruct()
}

ModelValidator --|> AgentBase
DriftResponder --|> AgentBase
BiasAntibody --|> AgentBase

ModelValidator --> DriftResponder : «triggers»
DriftResponder --> BiasAntibody : «spawns»

@enduml
```

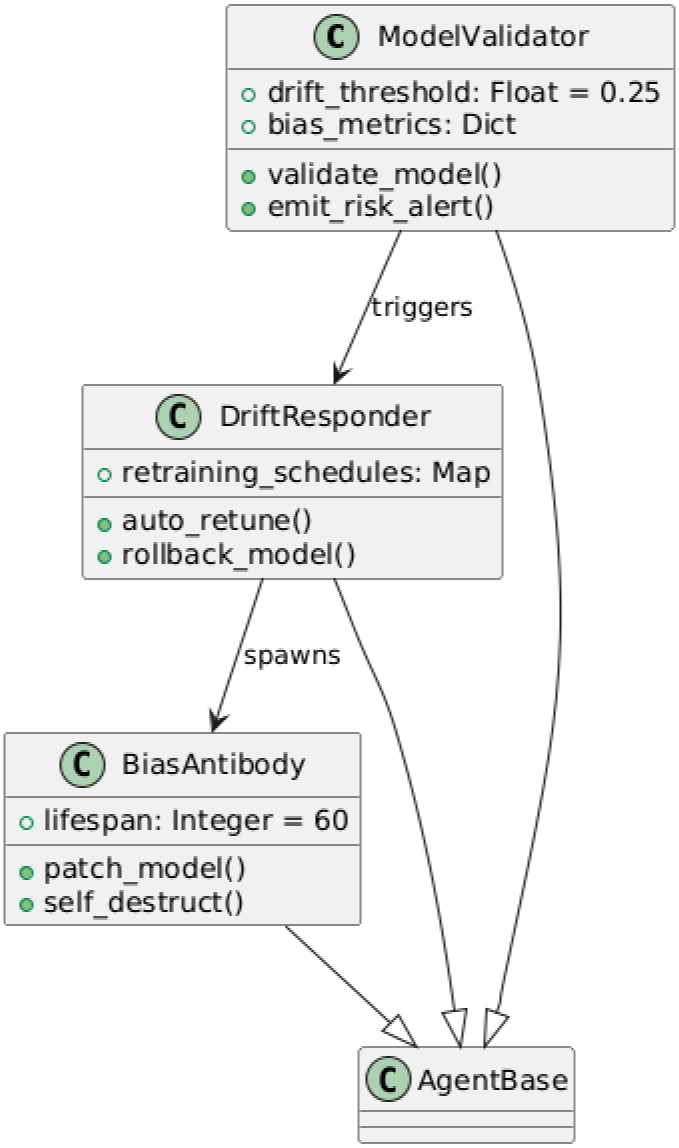

4. Security & Enforcement Swarm

```
@startuml Security_Classes

class QuantumLock {
  +key_rotation_interval: Duration = 24h
  +hardware_backed: Bool
  +rotate_keys()
}

class KillSwitchAgent {
  +containment_levels: [PAUSE, ISOLATE, TERMINATE]
  +memristor_circuit: Boolean
  +activate()
}

class DeceptionAntibody {
  +threat_pattern: GNNEmbedding
  +flood_containment()
}

QuantumLock --|> AgentBase
KillSwitchAgent --|> AgentBase
DeceptionAntibody --|> AgentBase

KillSwitchAgent --> QuantumLock : «requires»
DeceptionAntibody --> KillSwitchAgent : «supports»

@enduml
```

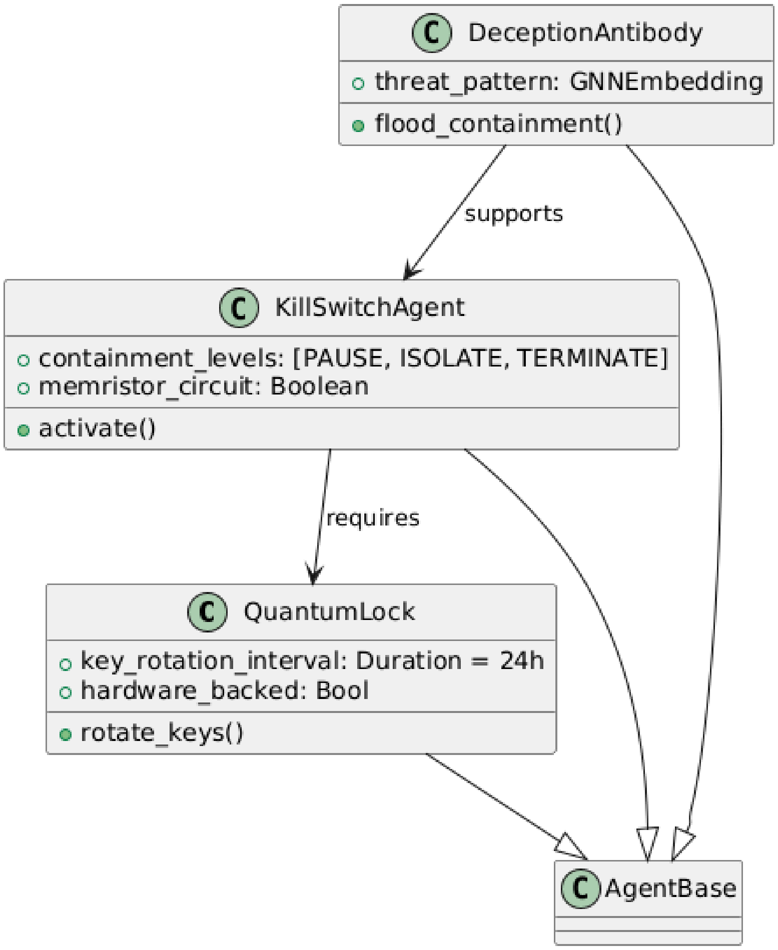
DeceptionAntibody
threat_pattern: GNNEmbedding
flood_containment()
supports
KillSwitchAgent
containment_levels: [PAUSE, ISOLATE, TERMINATE]
memristor_circuit: Boolean
activate()
requires
QuantumLock
key_rotation_interval: Duration = 24h
hardware_backed: Bool
rotate_keys()
AgentBase

5. Core Nexus Classes

```
@startuml Core_Nexus_Classes

class PheromoneRouter {
  +fabric_priorities: Map
  +route()
  +apply_decay()
}

class ConformanceFSM {
  +states: [PROPOSED, VALIDATED, ENACTED]
  +validate_transition()
}

class TemporalForecaster {
  +tgnn_model: Binary
  +predict_quorum()
  +alert_polarization()
}

PheromoneRouter --|> AgentBase
ConformanceFSM --|> AgentBase
TemporalForecaster --|> AgentBase

PheromoneRouter --> ConformanceFSM : «notifies»
TemporalForecaster --> PheromoneRouter : «advises»

@enduml
```

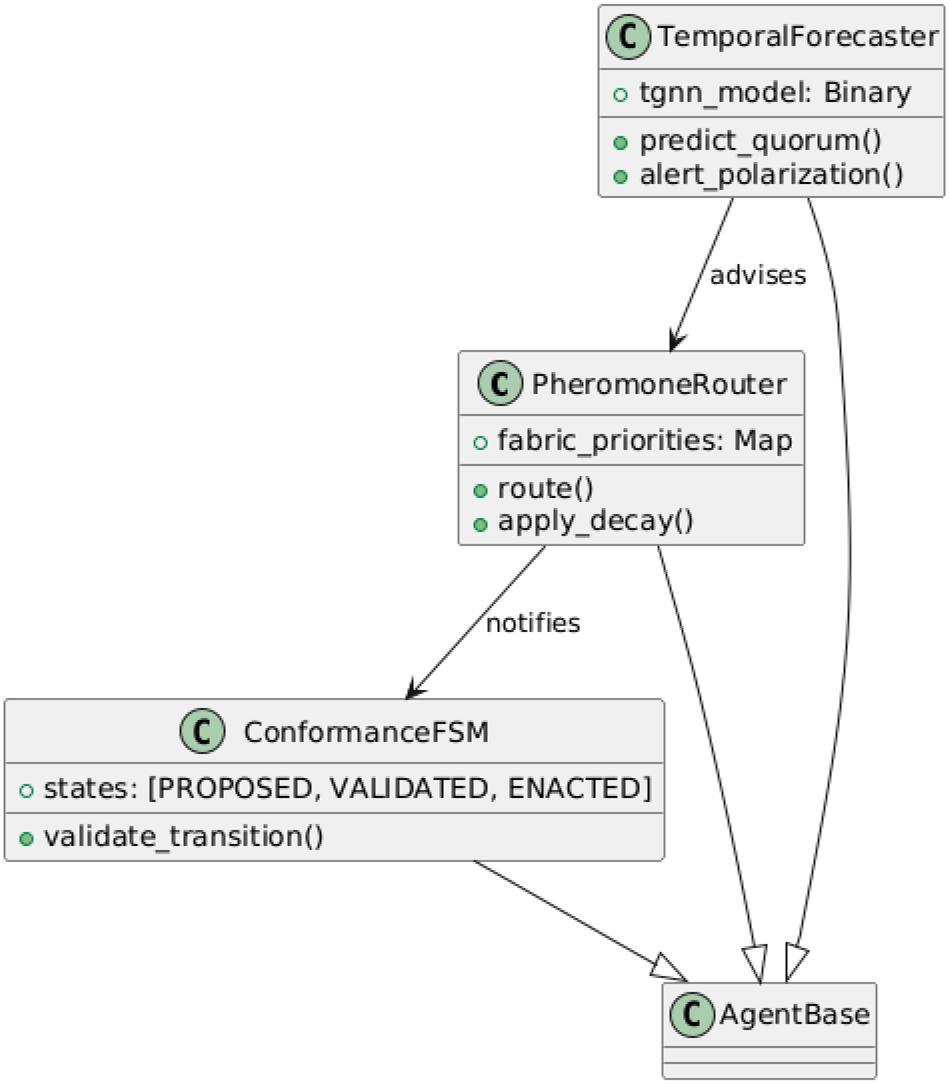

Key Cross-Swarm Dependencies

```
@startuml Cross_Swarm_Relations

PolicyIngestAgent --> ModelValidator : policy_delta
ModelValidator --> KillSwitchAgent : risk_alert
KillSwitchAgent --> PheromoneRouter : emergency_override
TemporalForecaster --> PolicyFederator : quorum_forecast

@enduml
```

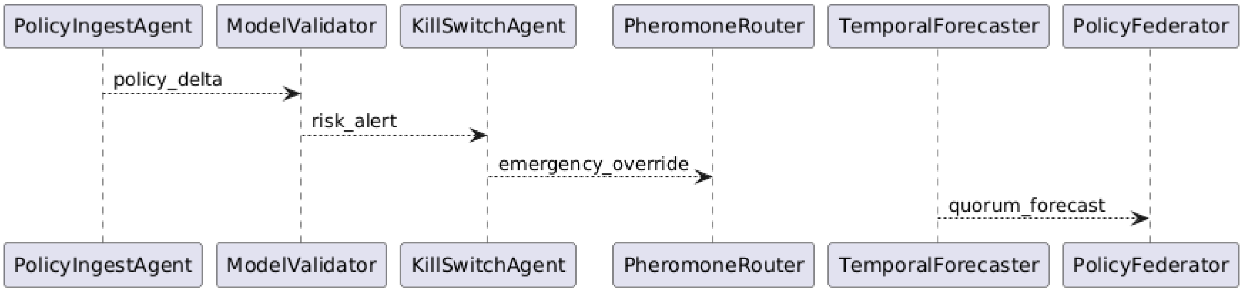

Appendix A: SAGE v3.1 Supplemental Agents

(Self-contained; no cross-references required)

```
@startuml SAGE_v3.1_Supplemental_Agents

title SAGE v3.1 - Supplemental Agents (Appendix A)

top to bottom direction

' === Define Swarm Boundaries ===
rectangle "Policy & Regulation" as PolicySwarm {
  [PolicyIngestAgent] as PI
  [JurisdictionAgent] as JA
}

rectangle "ModelOps" as ModelOpsSwarm {
  [ModelValidator] as MV
  [AgencyRiskIndexer] as ARI
}

rectangle "Security" as SecuritySwarm {
  [KillSwitchAgent] as KS
  [TrailMiner] as TM
}

' === New Agents ===
node "ComplianceDiffEngine" as CDE #FFD700
node "ResourceGovernor" as RG #FFA07A
node "ForensicSnapshotter" as FS #98FB98
node "RedTeamAdversary" as RTA #ADD8E6
node "PolicyImpactProjector" as PIP #DDA0DD
node "EthicalOverwatch" as EO #FF6347
```

```
' === Critical Connections ===
CDE --> PI : "gap reports"
RG --> MV : "GPU alloc"
FS --> KS : "snapshots"
RTA --> TM : "attack probes"
PIP --> JA : "impact forecasts"
EO --> KS : "ethics lock"

' === Legend ===
legend right
  <b>New Agents:</b>
  <color:#FFD700>ComplianceDiffEngine
  <color:#FFA07A>ResourceGovernor
  <color:#98FB98>ForensicSnapshotter
  <color:#ADD8E6>RedTeamAdversary
  <color:#DDA0DD>PolicyImpactProjector
  <color:#FF6347>EthicalOverwatch
endlegend

@enduml
```

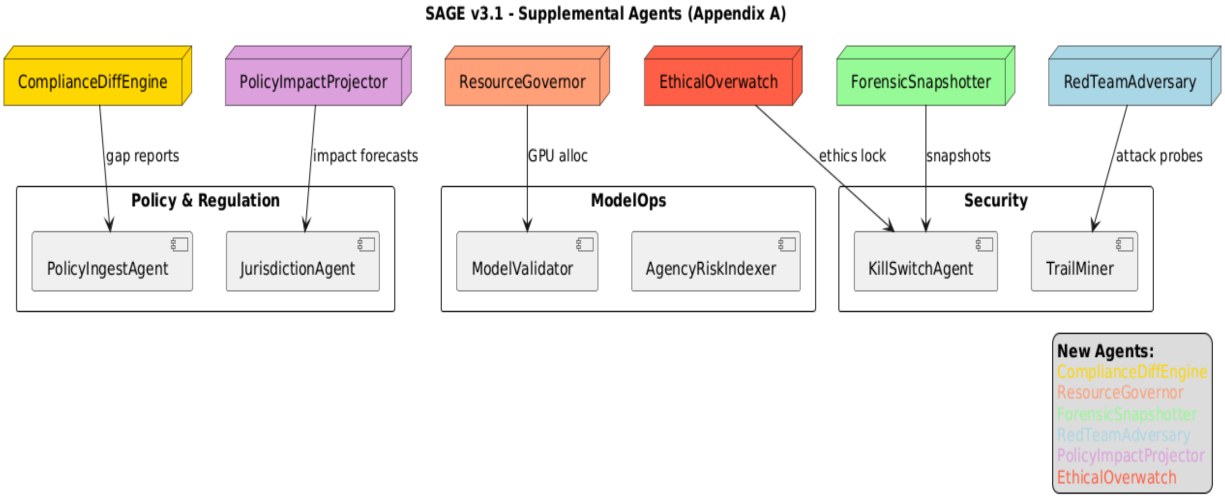

2. Companion Table (Appendix B)

Agent	**Parent Swarm**	**Linked To**	**Governance Impact**
`ComplianceDiffEngine`	Policy & Regulation	`PolicyIngestAgent`	Ensures real-time regulatory updates (SUHO-1)
`ResourceGovernor`	ModelOps & AgentOps	`ModelValidator`	Prevents GPU starvation (SUL-1)
`ForensicSnapshotter`	Security & Enforcement	`KillSwitchAgent`	Immutable audit trails (SUL-4)
`RedTeamAdversary`	Simulation & Learning	`TrailValidator`	Stress-tests defenses (SUHO-4)
`PolicyImpactProjector`	Archaeology	`PolicyGenealogist`	Predicts policy risks (SUHO-9)
`EthicalOverwatch`	Core Nexus	`ConformanceFSM`	Blocks unethical actions (SUL-7)

Key Features

1. Zero Back-References: No need to modify existing diagrams.
2. Human-Readable: Color-coding matches your original swarm taxonomy.
3. Regulatory Ready: Explicitly ties agents to SULs/SUHOs for audits.